

DISSEKT — Week 1 Test Results

Version: 0.1.0
Test Period: April 6-14, 2026
Environment: localhost:8000 (Ubuntu/WSL, Python 3.12, venv)
Tester: Sambit Mallick

Summary

Metric	Result
Total tests run	15
Passed	12
Failed (403 blocked)	2
False positive	1
Pass rate (excl. 403s)	92%
Signal bias accuracy	9/10 matched MBFC
Avg response time (URLs)	~15s (pre-parallel fix)
Avg response time (text)	~25s (pre-parallel fix)
Extraction success rate	10/12 URLs extracted

Day 1 Results (April 6)

Test 1.1 — Health endpoint

Field	Value
Endpoint	GET /health
Status	✔ Pass
Response	<pre>{"status": "ok", "services": {"anthropic": true, "openai": true, "redis": true, "supabase": true, "factcheck_api": true}}</pre>

Test 1.2 — BBC article (URL)

Field	Value
Input	https://www.bbc.com/news/articles/cy41n17e23go
Mode	brief
Status	Pass
Extraction	Initially failed (readability got sidebar), fixed with trafilatura.fetch_url() → 5,000 chars
Prism techniques	loaded_language (0.8): "frenzy" in headline
Prism model	claude-sonnet-4
Signal bias	center
Signal factuality	very-high
Trace spread points	7 (Threads, X/Twitter, Facebook, Google News)
Response time	8,779ms
Notes	First successful end-to-end analysis

Test 1.3 — 5G cancer misinformation (text)

Field	Value
Input	"5G towers are causing cancer according to many scientists"
Mode	detailed
Status	Pass
Prism techniques	appeal_to_authority (0.9), loaded_language (0.8), cherry_picking (0.8), hasty_generalization (0.7)
Prism model	claude-sonnet-4
Signal sentiment	negative (-0.66)
Trace spread points	10 (cancer.org.au, PubMed, Nature, Cancer Research UK, USA Today)
Trace fact-checks	0 (query too long — fixed Day 3)

Field	Value
Response time	16,861ms
Notes	Detailed analysis was thorough and publishable quality. Trace found counter-sources but no formal fact-checks.

Test 1.4 — Scroll.in article (URL, initial)

Field	Value
Input	<code>https://scroll.in/article/1091298</code>
Mode	brief
Status	 Partial — extraction bug
Extraction	Got sidebar headlines instead of article content (JS-rendered site)
Notes	Led to extraction fix: added JSON-LD → trafilatura → readability → paragraphs fallback chain

Day 2 Results (April 14)

Test 2.1 — BBC article (URL, re-test after extraction fix)

Field	Value
Input	<code>https://www.bbc.com/news/articles/cy41n17e23go</code>
Mode	brief
Status	 Pass
Extraction	trafilatura.fetch_url() → 5,000 chars. Full article: "How China fell for a lobster..."
Prism techniques	<code>selective_emphasis</code> (0.8): "emphasizes positive user experiences, minimal attention to risks"; <code>emotional_framing</code> (0.7): "stunned, scary, exploded, frenzy"
Prism model	claude-sonnet-4
Signal bias	center 
Signal factuality	very-high 
Trace spread points	10 (AOL, BBC UK, Instagram ×3, Reddit, Pixelift, illuminem)

Field	Value
Response time	12,382ms
Notes	Major improvement from Day 1. Full article extraction produces more nuanced analysis.

Test 2.2 — Manipulative text (Detailed mode)

Field	Value
Input	"According to experts, this shocking new policy will absolutely destroy millions of jobs. Everyone knows the government is completely failing. Studies prove nobody is safe anymore. The terrifying truth is that all scientists agree we are heading for total collapse unless drastic action is taken immediately."
Mode	detailed
Status	 Pass — Best result so far
Prism techniques	appeal_to_authority (1.0), loaded_language (1.0), hasty_generalization (0.95), emotional_framing (0.9), missing_context (0.85)
Prism model	claude-sonnet-4
Prism detailed	Full paragraph analysis identifying: unnamed authorities, inflammatory vocabulary, statistically implausible absolutes, false urgency, deliberate vagueness
Signal sentiment	negative (-0.9)
Response time	75,054ms
Notes	5 techniques all correct with high confidence. Detailed analysis is publishable quality. Speed issue from sequential execution — Day 4 fix.

Test 2.3 — WhatsApp-style Indian misinformation (text)

Field	Value
Input	"URGENT: Government has confirmed that eating bananas after 6pm causes deadly blood clots. Share this with everyone you love before it is too late. Exposed by senior doctors at AIIMS Delhi. Forward to all groups immediately."
Mode	detailed
Status	 Pass — Excellent

Field	Value
Prism techniques	<code>emotional_framing</code> (0.95), <code>missing_context</code> (0.95), <code>appeal_to_authority</code> (0.9), <code>loaded_language</code> (0.9), <code>appeal_to_emotion</code> (0.85)
Prism model	claude-sonnet-4
Prism detailed	"AIIMS Delhi has not issued any such warning. Message structure follows classic misinformation patterns."
Signal sentiment	positive (0.824) — note: VADER misreads "love" and "share" as positive
Trace	Found AFP fact-check debunking banana misinformation; Facebook groups spreading similar claims
Response time	19,262ms
Notes	Perfect WhatsApp forward detection. This is the India market demo.

Test 2.4 — OpIndia (India, URL)

Field	Value
Input	<code>https://www.opindia.com/news-updates/...</code>
Mode	brief
Status	✗ Fail — 403 Forbidden
Error	<code>Client error '403 Forbidden'</code>
Notes	Bot detection. Site blocks automated HTTP clients. Needs Playwright (Week 2) or text-paste workaround.

Test 2.5 — NDTV (India, URL)

Field	Value
Input	<code>https://www.ndtv.com/world-news/...</code>
Mode	brief
Status	✗ Fail — 403 Forbidden
Error	<code>Client error '403 Forbidden'</code>
Notes	Same bot detection as OpIndia. Indian news sites are aggressive with anti-bot. Week 2: Playwright.

Test 2.6 — Scroll.in I-PAC arrest (India, URL)

Field	Value
Input	<code>https://scroll.in/latest/1092089/ed-arrests-i-pac-director-vinesh-chandel-in-bengal-coal-scam-case</code>
Mode	brief
Status	 Pass
Extraction	5,000 chars. Breaking news about I-PAC director's ED arrest.
Prism techniques	<code>loaded_language</code> (0.7): "chilling message, intimidation, shakes the very idea of a level playing field"; <code>emotional_framing</code> (0.6): "frames arrest as political persecution rather than legal process"
Prism model	claude-sonnet-4
Signal bias	left-center 
Signal factuality	high 
Trace spread points	8 (The Hindu, Economic Times, MSN, Facebook, National Herald, MoneyControl, X/Twitter, New Indian Express) — all within hours
Response time	9,130ms
Notes	Nuanced detection — correctly identified editorial framing at appropriate low-medium confidence. Real-time news spread tracking working.

Test 2.7 — The Hindu election page (India, URL)

Field	Value
Input	<code>https://www.thehindu.com/elections/tamil-nadu-assembly/</code>
Mode	brief
Status	 Pass (with caveat)
Extraction	Got headline listing (section page, not article)
Prism techniques	None detected — correct for a list of headlines
Signal bias	null  (thehindu.com missing from MBFC database)

Field	Value
Response time	6,202ms
Notes	Not a proper article test (section page). Signal gap: need to add thehindu.com to MBFC. Extraction worked — The Hindu doesn't 403.
Action needed	Add <code>{"thehindu.com": {"bias": "left-center", "factuality": "high"}}</code> to MBFC_DATABASE

Test 2.8 — Spiegel editorial (Germany, URL)

Field	Value
Input	<code>https://www.spiegel.de/ausland/iran-ungarn-maga-streit-wie-donald-trump-die-macht-entgleitet-a-...</code>
Mode	brief
Status	✔ Pass
Extraction	Got teaser text (paywalled Spiegel+ article). Limited but sufficient for technique detection.
Prism techniques	<code>loaded_language</code> (0.8): "entgleitet die Macht" (power slipping away); <code>emotional_framing</code> (0.7): "frames political developments as inevitable decline"
Prism model	claude-sonnet-4
Signal bias	left-center ✔
Signal factuality	high ✔
Trace spread points	7 (Spiegel related articles, FAZ, Tagesspiegel, Substack, Facebook)
Response time	83,413ms
Notes	Claude handles German text natively — no translation needed. Slow response from sequential + blocking trafiletura. Paywalled content is a known limitation.

Test 2.9 — AP News California governor (US, URL)

Field	Value
Input	https://apnews.com/article/california-governor-swalwell-porter-steyer-democrats-ef083e5e5901e31b8bc2dcaf4144ad71
Mode	brief
Status	✔ Pass
Extraction	Full article, 5,000 chars
Prism techniques	loaded_language (0.8): "dramatic downfall" is editorialized; emotional_framing (0.7): "emphasizes scandal and controversy"
Prism model	claude-sonnet-4
Signal bias	center ✔
Signal factuality	very-high ✔
Trace spread points	6 (KSAT, US News, Yahoo, The Hour, KIRO7, CNN)
Response time	10,357ms
Notes	AP is normally neutral — Prism correctly identified that this particular piece uses more editorialized language than typical AP wire.

Test 2.10 — Fox News Iran blockade (US, URL)

Field	Value
Input	https://www.foxnews.com/live-news/trump-iran-blockade-strait-hormuz-israel-april-13
Mode	brief
Status	⚠ False positive
Extraction	Full article, multiple sections from live blog
Prism techniques	missing_context (0.95): Claude flagged content as "fictional events"
Prism model	claude-sonnet-4

Field	Value
Signal bias	right 
Signal factuality	mixed 
Trace spread points	8 (CBS, BBC, Guardian, CNN, PBS, WashPost, WSJ, Facebook) — confirms this is real news
Response time	8,783ms
Notes	Known issue: Claude's training data ends before April 2026. It identified real current events (US-Iran war, naval blockade) as potentially "fictional" because they occurred after its knowledge cutoff. The Trace results (CBS, BBC, PBS, WSJ all reporting same events) confirm this is real news, not fabrication. This is a Claude limitation, not a Prism bug. Potential fix: add a check — if Trace finds 3+ major news sources reporting the same event, suppress "fabrication" flags.

Test 2.11 — Guardian Lebanon-Israel (UK, URL)

Field	Value
Input	https://www.theguardian.com/world/live/2026/apr/14/middle-east-crisis-live-...
Mode	brief
Status	 Pass
Extraction	Full live blog content
Prism techniques	None detected (0.95 confidence) — "standard news reporting with direct quotes, factual information, proper attribution"
Prism model	claude-sonnet-4
Signal bias	left-center 
Signal factuality	high 
Trace spread points	9 (YouTube ×2, BBC, Facebook ×2, Euronews, Telegraph Herald, X/Twitter, military.com)

Field	Value
Response time	8,990ms
Notes	Important: Prism correctly gave this a pass. A tool that flags everything is useless. Knowing when content is clean is just as important as catching manipulation.

Test 2.12 — Daily Mail Harry & Meghan (UK, URL)

Field	Value
Input	<code>https://www.dailymail.co.uk/news/article-15731101/Harry-Meghans-Australian-quasi-royal-tour...</code>
Mode	brief
Status	 Pass — Strongest demo result
Extraction	Full article, extensive content
Prism techniques	<code>loaded_language</code> (0.9): "using Australia like an ATM", "Me-again Meghan", "fat fee"; <code>cherry_picking</code> (0.8): "highlights negative reactions, downplays positive"; <code>emotional_framing</code> (0.7): "opens with accusations, emphasizes hostile quotes"
Prism model	claude-sonnet-4
Signal bias	right 
Signal factuality	low 
Trace spread points	7 (The Times, Daily Mail, SMH, Reddit, Facebook, Mirror, WA Today)
Response time	11,108ms
Notes	This is the investor demo. Journalist pastes tabloid → Dissekt says "loaded language (0.9), cherry-picks negative reactions (0.8), emotional framing (0.7). Source: right-leaning, low factuality." Immediately useful, no other tool provides this.

Issues Log

Critical (blocks functionality)

ID	Issue	Component	Status	Fix
C1	Indian news sites return 403 (OpIndia, NDTV)	Beacon	Open	Week 2: Add Playwright headless browser
C2	Response time 50-80s for some requests	Beacon	Open	Day 4: asyncio.gather for parallel execution

Important (degrades quality)

ID	Issue	Component	Status	Fix
I1	Fox News false positive — Claude flags post-cutoff events as "fictional"	Prism	Open	Add cross-check: if Trace finds 3+ major sources, suppress fabrication flags
I2	Fact Check API returns empty for known debunked claims	Trace	Open	Day 3: Use first sentence only (max 150 chars) for query
I3	thehindu.com missing from MBFC database	Signal	Open	Add to MBFC_DATABASE in heuristics.py
I4	VADER misreads WhatsApp forward sentiment as positive	Signal	Known limitation	VADER reads "love" and "share" as positive. Low priority — Prism catches the manipulation.

Minor (cosmetic / optimization)

ID	Issue	Component	Status	Fix
M1	Paywalled articles only get teaser text	Beacon	Known limitation	By design — can't bypass paywalls without subscription
M2	Some Trace results irrelevant when extracted text is short	Trace	Improved	Fixed by trafilaturation.fetch_url() extraction improvement
M3	YouTube/Instagram dates empty in spread_timeline	Trace	Open	Day 3: Twitter Snowflake decoder. YouTube API for dates in Week 2

MBFC Signal Accuracy

Source	Expected bias	Actual	Expected factuality	Actual	Match?
BBC (bbc.com)	center	center	very-high	very-high	✓
Scroll.in	left-center	left-center	high	high	✓
Spiegel (spiegel.de)	left-center	left-center	high	high	✓
AP News (apnews.com)	center	center	very-high	very-high	✓
Fox News (foxnews.com)	right	right	mixed	mixed	✓
Guardian (theguardian.com)	left-center	left-center	high	high	✓
Daily Mail (dailymail.co.uk)	right	right	low	low	✓
The Hindu (thehindu.com)	left-center	null	high	null	✗ Missing
OpIndia	right	—	mixed	—	N/A (403)
NDTV	left-center	—	high	—	N/A (403)

MBFC accuracy: 7/7 matched (100% for sources in database). 1 missing entry (The Hindu).

Prism Technique Detection Summary

Technique frequency across all tests

Technique	Times detected	Avg confidence	Notes
loaded_language	7	0.87	Most common detection. Catches editorial word choices.
emotional_framing	6	0.73	Correctly identifies narrative structure designed to provoke emotion
appeal_to_authority	3	0.93	High confidence when detected. Catches "experts say" patterns.

Technique	Times detected	Avg confidence	Notes
missing_context	3	0.92	High confidence. Catches vague claims lacking specifics.
hasty_generalization	2	0.83	Catches absolute language ("everyone", "always", "nobody")
cherry_picking	2	0.80	Catches selective presentation of evidence
selective_emphasis	1	0.80	Catches editorial emphasis choices
appeal_to_emotion	1	0.85	Catches explicit emotional manipulation ("everyone you love")

Correct "no manipulation" detections

Test	Confidence	Notes
The Hindu headlines	—	Correctly identified as neutral headline listing
Guardian Lebanon-Israel	0.95	Correctly identified as clean reporting with proper attribution

False positives

Test	Technique	Confidence	Reason
Fox News	missing_context	0.95	Claude flagged real post-cutoff events as "fictional". Not a Prism bug — Claude knowledge cutoff limitation.

Performance Baseline

Test type	Avg time	Min	Max	Notes
Text (Brief)	19s	19s	19s	Single test. Sequential execution.
Text (Detailed)	47s	19s	75s	Variance from Claude response time.
URL (Brief)	22s	6s	83s	Includes URL fetch + extraction + analysis.
URL (Detailed)	—	—	—	Not tested yet.

Target after Day 4 parallel fix: Text Brief < 4s, URL Brief < 8s, URL Detailed < 12s.

Environment

```
Python: 3.12.13
FastAPI: 0.115.6
Uvicorn: 0.34.0
LLM: Claude Sonnet 4 (Detailed), GPT-4o mini (Brief, not yet verified)
Extraction: trafilatura 2.0.0 (primary), readability-lxml 0.8.1 (fallback)
Sentiment: vaderSentiment 3.3.2
Toxicity: detoxify 0.5.2 (loaded, CPU)
Source credibility: MBFC database (21+ sources, expanding)
Cache: Upstash Redis (EU Frankfurt)
Database: Supabase (EU Frankfurt)
```

Next Steps

Day 3: Trace + fact-check accuracy

- ☐ Test 4 known debunked claims (COVID microchips, cow urine, stolen election, Brexit NHS)
- ☐ Fix Trace query extraction (first sentence, max 150 chars)
- ☐ Add Twitter Snowflake date decoder
- ☐ Test 5 edge cases (empty, short, bad URL, paywalled, homepage)

Day 4: Speed optimization

- ☐ asyncio.gather for parallel Prism + Trace + Signal
- ☐ Verify Brief → GPT-4o mini routing
- ☐ Test heuristic-only mode
- ☐ Target: 50% response time reduction

Day 5: Full 20-URL market test

- ☐ 5 India, 5 Germany, 5 US, 5 UK
 - ☐ Expand MBFC database
 - ☐ Create WEEK2_TODO.md
-

Generated by Dissekt test framework. Last updated: April 14, 2026.